

Departmental Coversheet
Hillary term 2022 Mini-project
Paper title: Database System Implementation
Candidate Number: 1058016
Your degree: MSc Advanced Computer Science

To run the code, try make all. There will be an output file RDF.DB.out. Run the file and you can enter the database interface.

1. (a) The data structure in this paper is a six-column triple table. The subject, predicate, and object of a triple are encoded as integers in the first three columns, R_s , R_p , and R_o , respectively. Conceptually, there are three linked lists, an sp -list that connects all triples with the same R_s grouped by R_p , an op -list that associates all triple with the same R_o grouped by R_p , and a p -list that relates all triples with the same R_p without any grouping. In the table, the last three columns, N_{sp} , N_{op} , and N_p , store the next-pointes, which are going to be the row numbers in the triple table in the actual implementation.

Apart from the table, six index maps are also maintained. I_s , I_p , and I_o store the head of the sp -list, p -list, and op -list, respectively. I_{sp} maps s and p to the first occurrence of the triple with the same s and p in the sp -list. So does I_{op} . I_{spo} stores the row number of each triple in the table.

- (b) ADD(Triple t) consists of two parts; first, append a new row to the end of the RDF-index table, and second, associates the new row with all six index maps and alters the pointer columns, N_{sp} , N_{op} , and N_p , to point to the correct row.

Since there is no need to worry about the concurrency in our setting, the RDF-index table is maintained as a fixed size vector of int list of size 6 with each position stands for different columns. Therefore, for each time we add a triple into the table, we simply append it to the last. However, if a triple is already in the I_{spo} map, we skip it. When the size reaches its limit, the entire table will resize.

Algorithm 1 ADD (t)

```

if  $t$  in  $I_{spo}$  then                                     ▷  $t = (s, p, o)$  is a triple.
    return
end if
 $i$  = # of elements in the triple-table
if  $i + 1 >$  the size of the triple-table then
    resize the triple table
end if
 $T_{new} = [t.s, t.p, t.o, -1, -1, -1]$                      ▷ The last three columns are left for update later.
triple-table[ $i$ ] =  $T_{new}$ 
 $I_{spo}[t] = i$ 
Update remaining indexes

```

The columns, N_{sp} , N_{op} , and N_p , are maintained in a linked-list-like manner and are updated simultaneously with the index maps. Take N_{op} for example, if a new triple, T_{new} , does not appear in the I_o and I_{op} , it means that T_{new} is a triple with a brand new $T_{new}.s$ and $T_{new}.p$. Therefore, we insert T_{new} into the index map I_o and I_{op} (case1). If we found a T with $T.o = T_{new}.o$ and $T.p = T_{new}.p$, then we insert T_{new} after T , and point the next of T_{new} to the original next of the T (case3). A special case is that when there is no next for T and it is handled in case 2 of the Algorithm (2). The case of N_{sp} and N_p is similar.

Algorithm 2 Update $I_{op}(T_{new})$

```
 $T$  = the first triple with  $T.o = T_{new}.o$  and  $T.p = T_{new}.p$   
if  $T$  does not exist then ▷ Case 1  
    make  $T_{new}$  the head of  $I_o$  and  $I_{op}$   
end if  
if  $T$  does not have  $T_{next}$  then ▷ Case 2  
    make  $T_{new}$  the head of  $I_o$  and  $I_{op}$   
     $T_{new}.N_{op} = T$   
end if  
if  $T$  has  $T_{next}$  then ▷ Case 3  
     $T_{next} = T.N_{op}$   
     $T_{new}.N_{op} = T_{next}$   
end if
```

- (c) Let $t = \langle s, p, o \rangle$ be a triple pattern and X, Y , and Z be free variables. There are two inputs for EVALUATE; t , the matching pattern, and $index$, the previous row number. There are also two outputs, t' , the answer to the match pattern, $index$, the current row number. The index are usually the pointer (row number) to the next triple, but there are two sentinels; *EndOfNode* will tell the caller there is no more next triple to search and *EndSearch* stands for an illegal search. The general framework is shown in Algorithm (3).

Algorithm 3 Framework

```
while  $index \neq EndOfNode$  and  $index \neq EndSearch$  do  
     $index, t' \leftarrow Evaluate(t, index)$   
end while
```

Inside EVALUATE, there are four categories, from no free variable to three variables.

When there is not a free variable, EVALUATE will check if t is in I_{spo} map, if it is return the triple and set the index to be *EndOfNode*. Otherwise, it will return nothing and set the index be *EndSearch*. The pseudo-code is shown in Algorithm (4).

Algorithm 4 Evaluate $\langle s, p, o \rangle$

```
Require:  $t = \langle s, p, o \rangle$   
if  $t$  in  $I_{spo}$  then  
     $index \leftarrow EndOfNode, t' \leftarrow t$   
else  
     $index \leftarrow EndSearch, t' \leftarrow Null$   
end if
```

When there is one free variable, there are three cases: $\langle X, p, o \rangle$, $\langle s, X, o \rangle$, and $\langle s, p, X \rangle$. The first and the last cases are similar. For $\langle X, p, o \rangle$, in the first search, we use the index map, I_{op} , to locate the head of the triple with predicate and object the same as o and p , then we return the first match and the next pointer at N_{op} . For all subsequent evaluates, we traverse N_{op} list until there is not a match. The pseudo-code is suggested in Algorithm (5). $\langle s, p, X \rangle$ is similar.

Algorithm 5 Evaluate $\langle X, p, o \rangle$

Require: $t = \langle X, p, o \rangle$
 if $\text{hash}(t.p, t.o)$ does not appear in I_{op} **then**
 $\text{index} \leftarrow \text{EndSearch}$
 end if
 if it is the first search **then** ▷ locate the head
 $t' \leftarrow$ the triple I_{op} points to, $\text{index} \leftarrow t'.N_{op}$.
 else ▷ traverse N_{op} list
 $t' \leftarrow t.N_{op}$, $\text{index} \leftarrow t'.N_{op}$
 end if
 if $t'.o$ and $t'.p$ are not we are looking for **then** ▷ the end of grouped by.
 $\text{index} \leftarrow \text{EndSearch}$
 end if

Since we don't have an index on s and o , $\langle s, X, o \rangle$ needs some extra works. To match this pattern, we can traverse either sp -list or op -list and skip the triples does not match on s and o . The choose of the list depends on the size of $I_s[s]$ and $I_o[o]$ and we use the smaller one. The pseudo-code is in Algorithm (6).

Algorithm 6 Evaluate $\langle s, X, o \rangle$

Require: $t = \langle s, X, o \rangle$, where X is a special code stands for a free variable.
 if $|I_s[s]| < |I_o[o]|$ **then**
 do
 if it is the first search **then**
 $t' \leftarrow$ the triple I_{sp} points to, $\text{index} \leftarrow t'.N_{sp}$.
 else
 $t' \leftarrow t.N_{sp}$, $\text{index} \leftarrow t'.N_{sp}$
 end if
 while $t'.s \neq s$ and $t'.o \neq o$ and $\text{index} \neq \text{EndOfNode}$
 else
 do
 if it is the first search **then**
 $t' \leftarrow$ the triple I_{op} points to, $\text{index} \leftarrow t'.N_{op}$.
 else
 $t' \leftarrow t.N_{op}$, $\text{index} \leftarrow t'.N_{op}$
 end if
 while $t'.s \neq s$ and $t'.o \neq o$ and $\text{index} \neq \text{EndOfNode}$
 end if
 if $t'.s \neq s$ and $t'.o \neq o$ **then** ▷ corner case
 $\text{index} = \text{EndSearch}$
 end if

When there are two free variables, there are also three cases: $\langle X, Y, o \rangle$, $\langle X, p, Z \rangle$, and $\langle s, Y, Z \rangle$, where X , Y , and Z can be equal. Their ideas are similar; $\langle X, Y, o \rangle$, for example, we first find $I_o[t.o]$ and traverse the op -list till the end. If $X = Y$, then we skip those triples with $p \neq o$. The pseudo-code is given in Algorithm (7)¹.

¹ Even I found myself a little confused by what I wrote in the do-while loop, so this footnote might be helpful. After a few minutes of thinking, I found only when we want $X = Y$, the condition $t'.p \neq t'.s$ will make the loop start and try to locate the one with $t.p = t.s$. If $X \neq Y$, the loop will never start and return immediately.

Algorithm 7 Evaluate $\langle X, Y, o \rangle$

Require: $t = \langle X, Y, o \rangle$, where X and Y are special codes stand for variables.

```
do
  if it is the first search then
     $t' \leftarrow$  the triple  $I_o$  points to,  $index \leftarrow t'.N_{op}$ .
  else
     $t' \leftarrow t.N_{op}$ ,  $index \leftarrow t'.N_{op}$ 
  end if
  while  $X = Y$  and  $t'.p \neq t'.s$  and  $index \neq EndOfNode$ 
  if  $X = Y$  and  $t'.p \neq t'.s$  then  $\triangleright$  stop traversing when  $t'.p \neq t'.s$  when the pattern asks  $X = Y$ 
     $index = EndSearch$ 
  end if
```

The ideas for the other two patterns are similar; start from I_p or I_s map and traverse the p -list or sp -list.

When there are three free variables, the paper suggests to match, for example, the patterns like $\langle X, Y, Z \rangle$, we need to iterate over the triple table; if we want $X = Y$, we skip those $X \neq Y$. However, this is not efficient. Therefore, I modified a little bit and the idea and pseudo-code are discussed in 2.a, where I put everything that is different from the paper there.

2. (a) RDF indexing data structure that implements Add and Evaluate functions.

This component is included in `RDF_index.cpp`, in which ADD and EVALUATE are implemented as suggested in the problem 1. There are a few things that is slightly different from the paper.

1) I used XXHASH² by Facebook instead of Jenkins hashing, because it achieves state-of-the-art excellent performance on both long and small inputs. Jenkins hashing is implemented in `RDF_index.h` and but it is commented. I did some tests and observed that Jenkins is about 4% slower than XXHASH in average³.

2) Instead of open addressing, I eventually choose `std::unordered_map` for index maps I_{sp} , I_{op} and I_{spo} . I had an open addressing hash implemented in `HashTable.cpp` and `HashTable.h` (attached in the submission), but it was much slower than the `unordered_map` by about 20% or more.

3) To match the patterns like $\langle X, Y, Z \rangle$, the paper suggests to iterate over the triple table; if we want $X = Y$, we skip those $X \neq Y$. I modify this a little bit to improve the efficiency. Again, $\langle X, X, Z \rangle$, for example, we first iterate I_s , and for each s in I_s , we find if I_{sp} includes `hash(s, s)`, if yes, we traverse over the triple table. As shown in the pseudo-code in Algorithm (8)⁴.

Algorithm 8 Evaluate $\langle X, X, Z \rangle$

```
for  $s$  in  $I_s$  do
   $i = \text{hash}(s, s)$ 
  if  $i$  in  $I_{sp}$  then
    Evaluate_SPZ( $s, s$ )
  end if
end for
```

$\langle X, Y, Y \rangle$ and $\langle X, Y, X \rangle$ are similar. For $\langle X, X, X \rangle$, we only need to traverse s in I_s and find if $\langle s, s, s \rangle$ is in I_{spo} as shown in Algorithm (9).

²<https://cyan4973.github.io/xxHash/>

³For loading, XXHASH and Jenkins took 314.359, 4026.82, and 323.914, 4206.85, and 58582.4 ms on three datasets, respectively, achieving 2%, 4%, and 8% speed up on each.

⁴The actual implementation is slightly different. I put the check of if i in I_{sp} at the beginning of the Evaluate_SPZ to reduce some code redundancy, but they are equivalent.

Algorithm 9 Evaluate $\langle X, X, X \rangle$

```
for  $s$  in  $I_s$  do  
    Evaluate_SPO( $s, s, s$ )  
end for
```

- (b) The engine for evaluating BGP SPARQL queries.

This part is in the file SPARQL_engine.cpp and it is implemented strictly based on the model answer. A few additional lines are added for printing or improving performance.

- (c) The greedy join order optimization query planner.

This part can be sought in the file query_planner.cpp. I see what the model answer is trying to say, but I believe there are some minor mistakes, so I made some slight changes, but it is still $O(n^2)$, where n is the number of triple patterns we try to fit. There were four things I modified.

Firstly, the last three lines should go to the outside while loop.

Secondly, I changed the criteria for updating the new triple patterns. In the model answer, the criteria is

$$t_{best} = \perp \text{ or } score < score_{best} \text{ and either } var(t) = \emptyset \text{ or } var(t) \cap B \neq \emptyset. \quad (1)$$

It is troublesome because it will always take the first unprocessed triple pattern as the best pattern and compare this with the remaining. However, it might cause some issues in many cases. For example, the following triple patterns

$$\begin{aligned} \langle X, 1, 2 \rangle \\ \langle Y, 2, 3 \rangle \\ \langle X, 4, Y \rangle \end{aligned}$$

will produce the plan $\langle X, 1, 2 \rangle \mapsto \langle Y, 2, 3 \rangle \mapsto \langle X, 4, Y \rangle$ because even after we replace the X in $\langle X, 4, Y \rangle$, this triple pattern still has a higher selectivity than $\langle Y, 2, 3 \rangle$. Therefore, this plan will create an unnecessary cross product, and I believe the criteria (1) meant choosing the triple pattern that has the lowest selectivity and includes some variables that are in the processed patterns. Thus, I made some modifications to the algorithm, and it will 1) as long as an unprocessed pattern has a variable that is processed, we will always consider this pattern and 2) if this pattern has lower selectivity than the previous one, we choose this pattern as the best triple. With this modification, the plan becomes $\langle X, 1, 2 \rangle \mapsto \langle X, 4, Y \rangle \mapsto \langle Y, 2, 3 \rangle$.

Thirdly, there are some cases that $|\langle X, 1, 2 \rangle| = 10,000$ but $|\langle Y, 2, 3 \rangle| = 10$. If we ignore the size of the pattern, we might produce a lot of unnecessary joins, but we can store this information while we are updating I_{op} . Therefore, I also record the size of $I_{op}[op]$ and $I_{sp}[sp]$, and we choose the best triple pattern based on both selectivity and size. Our final query plan is $\langle Y, 2, 3 \rangle \mapsto \langle X, 4, Y \rangle \mapsto \langle X, 1, 2 \rangle$, which is the most optimal plan. Notice that only I_{op} and I_{sp} are considered because they are the most common cases, and recording the size of other index maps would produce overhead in indexing. Since the size of $I_{op}[t]$ and $I_{sp}[t]$ can be recored while we are inserting triple, the cost of maintaining the size is negligible⁵.

Lastly, the original query plan was based on the following precedence relation \prec on selectivity of triple patterns:

$$(s, p, o) \prec (s, ?, o) \prec (?, p, o) \prec (s, p, ?) \prec (?, ?, o) \prec (s, ?, ?) \prec (?, p, ?) \prec (?, ?, ?).$$

Now, since we have maintained the size of $(?, p, o)$ and $(s, p, ?)$, a new order of the selectivity should be used:

$$(s, p, o) \prec (s, ?, o) \prec (?, p, o) = (s, p, ?) \prec (?, ?, o) \prec (s, ?, ?) \prec (?, p, ?) \prec (?, ?, ?).$$

⁵There are many tricks to save memory. For example, let's say I_{sp} is an unordered_map<int, int>, where the first int is the hash key and the second int is the index of row number. For a table size smaller than 2^{24} rows, we can use the higher 7 bits to store the size and lower 24 bits to store the index. Other trick might apply also.

Therefore, when we are precessing these two patterns, instead of comparing the selectivity, we can directly compare the actual size of the triple patterns and select the smaller one. For example, for the following triple patterns,

$$\begin{aligned} &\langle X, 1, 2 \rangle \\ &\langle 2, 3, Y \rangle \end{aligned}$$

with $|\langle 2, 3, Y \rangle| > |\langle X, 1, 2 \rangle|$ will produce the following query plan,

$$\begin{aligned} &\langle 2, 3, Y \rangle \\ &\langle X, 1, 2 \rangle \end{aligned}$$

The pseudo-code goes as the following.

Algorithm 10 New-Plan-Query(U)

```

 $P \leftarrow [], B \leftarrow \emptyset$ 
while  $U \neq \emptyset$  do
   $t_{best} \leftarrow \perp, score_{best} \leftarrow 100, intersected \leftarrow \text{false}$ 
  for each unprocessed tripple pattern  $t \in U$  do
     $score \leftarrow$  the position of  $t$  in  $\prec$  where the variables in  $B$  are considered bounded
    if  $t_{best} = \perp$  then
       $intersected = (var(t) \cap B \neq \emptyset)$ 
       $t_{best} \leftarrow t, score_{best} \leftarrow score$ 
    end if
    if  $t$  and  $t_{best}$  are of the form  $\langle X, p, o \rangle$  or  $\langle s, p, X \rangle$  then
       $t_{best} \leftarrow$  the one with a smaller size
    end if
    if  $intersected$  then
      if  $score < score_{best}$  and either  $var(t) = \emptyset$  or  $var(t) \cap B \neq \emptyset$  then
         $t_{best} \leftarrow t, score_{best} \leftarrow score$ 
      end if
    else
      if  $var(t) = \emptyset$  or  $var(t) \cap B \neq \emptyset$  then  $\triangleright$  make sure that the next pattern share some variables.
         $t_{best} \leftarrow t, score_{best} \leftarrow score, intersected \leftarrow \text{true}$ 
      end if
    end if
  end for
   $P \leftarrow P \cup t_{best}, B \leftarrow B \cup var(t_{best})$ 
end while

```

From Table 1, we see a significant improvement in most of the queries. The testing protocol is the same as the one in question 3. To do the test, uncomment lines 20-25 in query_planner.cpp and comment lines 26-53 and MAKE TEST. The analysis of the result is included in 3.c.

- (d) The component for parsing and importing Turtle files.
This part is included in Turtle_handler.cpp.
- (e) The parser for SPARQL queries
This part is included in query_parser.cpp.
- (f) The component implementing the command line
This part is included in interface.cpp.

	Model Answer (ms)	New Query Plan (ms)	Improvement (%)
q1	0.2269	0.0038	5971.05263
q2	370.891	0.8707	42596.8761
q3	1.289	0.0022	58590.9091
q4	0.1669	0.0241	692.53112
q5	2.2444	0.06	3740.66667
q6	0.3	0.2121	141.442716
q7	2716.41	0.0106	25626509.4
q8	28.3936	1.6634	1706.96164
q9	1033.41	1.1611	89002.6699
q10	2.3399	0.0018	129994.444
q11	0.0617	0.0205	300.97561
q12	0.0517	0.0105	492.380952
q13	1.3423	0.0018	74572.2222
q14	0.2384	0.1947	122.444787

Table 1: Time took to process and evaluate the query by the model answer and the new query plan in ms on the dataset LUBM-001-mat.ttl. The new plan achieves more than 1,859,602.501% speed up in average on the 14 queries.

3.a (a) Hardware and Software Configuration:

- i. Model Identifier: MacBookPro14,3
- ii. Processor Name: Apple M1 Pro
- iii. Processor Speed: 3.2 GHz
- iv. Number of Processors: 1
- v. Total Number of Cores: 10
- vi. Memory: 32 GB
- vii. Operating System: macOS Monterey, Version 12.3
- viii. Compiler Version: Apple clang version 13.1.6 (clang-1316.0.21.2)

- (b) We set a timeout if one instruction takes more than 3 minutes. The following protocol can be found in protocol.cpp. If you want to run the load test, fill in the correct path to the data and uncomment some lines. It will take about 15 minutes to run.

Loading Data Test Protocol:

- i. *Restart computer* and turn off all irrelevant applications.
- ii. Start a clean terminal.
- iii. Load dataset 001 and markdown the time. If timeout, mark TO.
- iv. Delete all objects.
- v. Load dataset 010 and markdown the time. If timeout, mark TO.
- vi. Delete all objects.
- vii. Load dataset 100 and markdown the time. If timeout, mark TO.
- viii. Delete all objects.
- ix. Go back to step ii repeat this process for 10 times.
- x. Take the average.

To eliminate the cache effect, I load 3 datasets one by one for 10 times instead of running a single dataset for 10 times. This will flush most of caches. Also, I should play my phone and don't touch my laptop while running.

- (c) We choose COUNT instead of SELECT and for each query. We set a timeout if one instruction takes more than 3 minutes.

Loading Data Test Protocol:

- i. *Restart computer* and turn off all irrelevant applications.
- ii. Start a clean terminal.
- iii. MAKE all and run the output file.
- iv. Load one dataset.
- v. Run 14 queries one by one and markdown the time.
- vi. Go back to the step v and repeat 10 times.
- vii. Take the average.

To avoid the cache effect, I run 14 queries one by one for 10 times instead of running a single queries for 10 times.

- 3.b The time needed to load and index the data and the time needed to produce all query answers for each RDF graph and query are showed in Table 2 and Table 3, respectively.

	LUBM-001-mat (ms)	LUBM-010-mat (ms)	LUBM-100-mat (ms)
Load & Index Time	114.651	2058.07	25624.2

Table 2: Time needed to load and index the data in ms.

	LUBM-001-mat (ms)	LUBM-010-mat (ms)	LUBM-100-mat (ms)
q1	0.0038	0.0063	0.0371
q2	0.8707	31.8538	348.732
q3	0.0022	0.008	0.012
q4	0.0241	0.0599	0.0932
q5	0.06	0.1908	0.2434
q6	0.2121	11.3648	253.683
q7	0.0106	0.0314	0.1369
q8	1.6634	4.153	9.8969
q9	1.1611	42.0941	981.896
q10	0.0018	0.0058	0.0163
q11	0.0205	0.0599	0.1568
q12	0.0105	0.2159	0.4706
q13	0.0018	0.01	0.2973
q14	0.1947	8.8916	272.475

Table 3: Time needed to produce all query answers in ms.

- 3.c q1: This query consists of two triple patterns, and they are both in the form of $\langle X, p, o \rangle$. My query plan picks the second one as a start because it is considerably smaller (1874 v.s. 4). Otherwise, it will be very slow, as Table 1 suggested.
- q2: This query contains two groups of the forms $\langle X, p, o \rangle$, and $\langle X, p, Z \rangle$, and each group has three triples. The plan is to find a $\langle X, p, o \rangle$, then a $\langle X, p, Z \rangle$ that matches the first pick, and pick whatever matches the variable, Z , in the remaining unprocessed patterns. It picks not only the least selectivity pattern but also the smallest pattern in each step, resulting in a very short running time (compared to the model solution in Table 1).

- q3: This one is similar to q1, but they have different sizes in each pattern. In LUBM-001-mat.ttl, q3 has 5999 possible matches in the first pattern and 6 in the second; q1 has 1874 in the first and 4 in the second. Therefore, q3 is slightly slower than q1.
- q4: Though this query looks terrifying, the speed is fast because all patterns share the same X in p , and besides this, there is at most one free variable in each. Therefore, once X with the smallest size is picked, the remaining is just to check if a Y can go with the X .
- q5: This one is similar to q1 and q3 but is much slower because the search space is large. This one has 8330 possibilities for the first pattern and 719 for the second, so even starting from the second one would slow down by a lot.
- q6: This query simply matches just one pattern. No plan is needed. From the smallest to the largest dataset, the result sizes are 7790, 99566, and 1048532. So the time increases as we expected.
- q7: There are two triples like $\langle X, p, o \rangle$. Since the first one has the size 7790 and the second one is 1627, starting at the second one, the performance is much better (see Table 1).
- q8: All three datasets produce 7790 results, so the time is about the same. It also illustrates that we have an excellent query plan because pattern 1 has 7790 matches in the smallest dataset but 1048532 matches in the largest dataset. However, pattern 3 has 239 possible matches among all datasets. Therefore, by choosing pattern 3 as a start, we ensure that the time does not change as the size of pattern 1 changes.
- q9: It is like q8, but the sizes of the first three patterns increase by 12x from the smallest dataset to the largest. Therefore, even if we pick the smallest pattern to start with, the time will still increase.
- q10: Same as q1 and q3.
- q11: Like q10, the performance of q1, q3, q10, and q11 are good because they choose the pattern with the smallest size to start.
- q12: It is almost like q11. It first finds the smallest of the form $\langle X, p, o \rangle$ and matches $\langle X, p, Y \rangle$. Then all others are just $\langle s, p, o \rangle$ and can be verified very quickly.
- q13: There are 8330 choices for the first pattern but just one for the second. If we do it in order, it will be slow. However, the plan first takes $\langle s, p, X \rangle$ and then $\langle X, p, o \rangle$ because they have the same selectivity but different in sizes.
- q14: Same as q6. The running time increases as the size increases.